

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats import kurtosis, skew
```

```
In [2]: df=pd.read_csv("C:/Users/USER/Desktop/Datasets/churn_Modelling.csv")
```

```
In [3]: df.tail()
```

```
Out[3]:
```

| | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender | Age | Tenure | Balance |
|------|-----------|------------|-----------|-------------|-----------|--------|-----|--------|----------|
| 9995 | 9996 | 15606229 | Obijaku | 771 | France | Male | 39 | 5 | 0.0 |
| 9996 | 9997 | 15569892 | Johnstone | 516 | France | Male | 35 | 10 | 57369.0 |
| 9997 | 9998 | 15584532 | Liu | 709 | France | Female | 36 | 7 | 0.0 |
| 9998 | 9999 | 15682355 | Sabbatini | 772 | Germany | Male | 42 | 3 | 75075.0 |
| 9999 | 10000 | 15628319 | Walker | 792 | France | Female | 28 | 4 | 130142.0 |

```
In [4]: df.dtypes
```

```
Out[4]: RowNumber      int64
CustomerId    int64
Surname        object
CreditScore    int64
Geography      object
Gender         object
Age            int64
Tenure         int64
Balance        float64
NumOfProducts  int64
HasCrCard      int64
IsActiveMember int64
EstimatedSalary float64
Exited         int64
dtype: object
```

```
In [5]: df.shape
```

```
Out[5]: (10000, 14)
```

In [6]: df

Out[6]:

| | RowNumber | CustomerId | Surname | CreditScore | Geography | Gender | Age | Tenure | Balance |
|------|-----------|------------|-----------|-------------|-----------|--------|-----|--------|----------|
| 0 | 1 | 15634602 | Hargrave | 619 | France | Female | 42 | 2 | 0.0 |
| 1 | 2 | 15647311 | Hill | 608 | Spain | Female | 41 | 1 | 83807.8 |
| 2 | 3 | 15619304 | Onio | 502 | France | Female | 42 | 8 | 159660.8 |
| 3 | 4 | 15701354 | Boni | 699 | France | Female | 39 | 1 | 0.0 |
| 4 | 5 | 15737888 | Mitchell | 850 | Spain | Female | 43 | 2 | 125510.8 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9995 | 9996 | 15606229 | Obijiaku | 771 | France | Male | 39 | 5 | 0.0 |
| 9996 | 9997 | 15569892 | Johnstone | 516 | France | Male | 35 | 10 | 57369.6 |
| 9997 | 9998 | 15584532 | Liu | 709 | France | Female | 36 | 7 | 0.0 |
| 9998 | 9999 | 15682355 | Sabbatini | 772 | Germany | Male | 42 | 3 | 75075.3 |
| 9999 | 10000 | 15628319 | Walker | 792 | France | Female | 28 | 4 | 130142.7 |

10000 rows × 14 columns



```
In [7]: del df["CustomerId"]
del df["Surname"]
```

In [8]: df

Out[8]:

| | RowNumber | CreditScore | Geography | Gender | Age | Tenure | Balance | NumOfProducts | HasCreditCard |
|------|-----------|-------------|-----------|--------|-----|--------|-----------|---------------|---------------|
| 0 | 1 | 619 | France | Female | 42 | 2 | 0.00 | 1 | |
| 1 | 2 | 608 | Spain | Female | 41 | 1 | 83807.86 | 1 | |
| 2 | 3 | 502 | France | Female | 42 | 8 | 159660.80 | 3 | |
| 3 | 4 | 699 | France | Female | 39 | 1 | 0.00 | 2 | |
| 4 | 5 | 850 | Spain | Female | 43 | 2 | 125510.82 | 1 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9995 | 9996 | 771 | France | Male | 39 | 5 | 0.00 | 2 | |
| 9996 | 9997 | 516 | France | Male | 35 | 10 | 57369.61 | 1 | |
| 9997 | 9998 | 709 | France | Female | 36 | 7 | 0.00 | 1 | |
| 9998 | 9999 | 772 | Germany | Male | 42 | 3 | 75075.31 | 2 | |
| 9999 | 10000 | 792 | France | Female | 28 | 4 | 130142.79 | 1 | |

10000 rows × 12 columns



```
In [9]: Idf=df.iloc[:,0:9]
        Odf=df.iloc[:, -1]
```

```
In [10]: Odf
```

```
Out[10]: 0      1
         1      0
         2      1
         3      0
         4      0
         ..
        9995    0
        9996    0
        9997    1
        9998    1
        9999    0
        Name: Exited, Length: 10000, dtype: int64
```

```
In [11]: Idf
```

```
Out[11]:
```

| | RowNumber | CreditScore | Geography | Gender | Age | Tenure | Balance | NumOfProducts | Has |
|------|-----------|-------------|-----------|--------|-----|--------|-----------|---------------|-----|
| 0 | 1 | 619 | France | Female | 42 | 2 | 0.00 | 1 | |
| 1 | 2 | 608 | Spain | Female | 41 | 1 | 83807.86 | 1 | |
| 2 | 3 | 502 | France | Female | 42 | 8 | 159660.80 | 3 | |
| 3 | 4 | 699 | France | Female | 39 | 1 | 0.00 | 2 | |
| 4 | 5 | 850 | Spain | Female | 43 | 2 | 125510.82 | 1 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... | |
| 9995 | 9996 | 771 | France | Male | 39 | 5 | 0.00 | 2 | |
| 9996 | 9997 | 516 | France | Male | 35 | 10 | 57369.61 | 1 | |
| 9997 | 9998 | 709 | France | Female | 36 | 7 | 0.00 | 1 | |
| 9998 | 9999 | 772 | Germany | Male | 42 | 3 | 75075.31 | 2 | |
| 9999 | 10000 | 792 | France | Female | 28 | 4 | 130142.79 | 1 | |

10000 rows × 9 columns



In [12]: `Idf.describe()`

Out[12]:

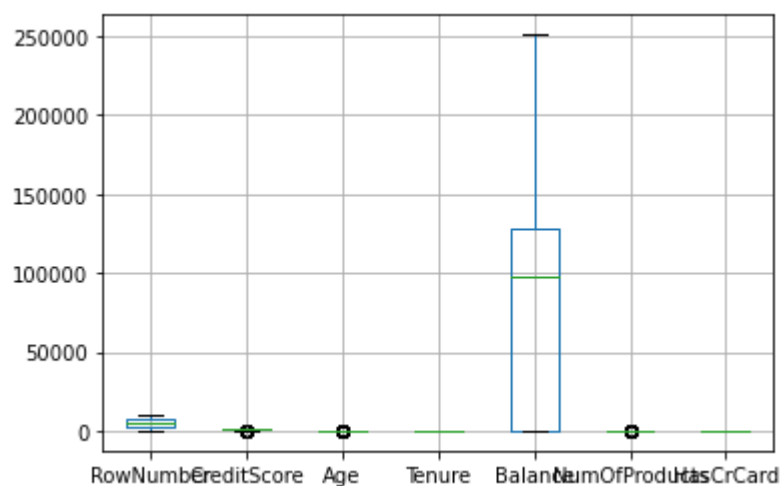
| | RowNumber | CreditScore | Age | Tenure | Balance | NumOfProducts | HasCrCard |
|--------------|-------------|--------------|--------------|--------------|---------------|---------------|--------------|
| count | 10000.00000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 | 10000.000000 |
| mean | 5000.50000 | 650.528800 | 38.921800 | 5.012800 | 76485.889288 | 1.530200 | 0.699510 |
| std | 2886.89568 | 96.653299 | 10.487806 | 2.892174 | 62397.405202 | 0.581654 | 0.460716 |
| min | 1.00000 | 350.000000 | 18.000000 | 0.000000 | 0.000000 | 1.000000 | 0.000000 |
| 25% | 2500.75000 | 584.000000 | 32.000000 | 3.000000 | 0.000000 | 1.000000 | 0.000000 |
| 50% | 5000.50000 | 652.000000 | 37.000000 | 5.000000 | 97198.540000 | 1.000000 | 0.699510 |
| 75% | 7500.25000 | 718.000000 | 44.000000 | 7.000000 | 127644.240000 | 2.000000 | 0.699510 |
| max | 10000.00000 | 850.000000 | 92.000000 | 10.000000 | 250898.090000 | 4.000000 | 1.000000 |

In [13]: `Idf.isnull().sum()`

Out[13]:

| | |
|---------------|-------|
| RowNumber | 0 |
| CreditScore | 0 |
| Geography | 0 |
| Gender | 0 |
| Age | 0 |
| Tenure | 0 |
| Balance | 0 |
| NumOfProducts | 0 |
| HasCrCard | 0 |
| dtype: | int64 |

In [14]: `boxplot=Idf.boxplot()`



```
In [15]: Newdf=pd.get_dummies(Idf)
Newdf
```

```
Out[15]:
```

| | RowNumber | CreditScore | Age | Tenure | Balance | NumOfProducts | HasCrCard | Geography_F |
|------|-----------|-------------|-----|--------|-----------|---------------|-----------|-------------|
| 0 | 1 | 619 | 42 | 2 | 0.00 | 1 | 1 | |
| 1 | 2 | 608 | 41 | 1 | 83807.86 | 1 | 0 | |
| 2 | 3 | 502 | 42 | 8 | 159660.80 | 3 | 1 | |
| 3 | 4 | 699 | 39 | 1 | 0.00 | 2 | 0 | |
| 4 | 5 | 850 | 43 | 2 | 125510.82 | 1 | 1 | |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 9995 | 9996 | 771 | 39 | 5 | 0.00 | 2 | 1 | |
| 9996 | 9997 | 516 | 35 | 10 | 57369.61 | 1 | 1 | |
| 9997 | 9998 | 709 | 36 | 7 | 0.00 | 1 | 0 | |
| 9998 | 9999 | 772 | 42 | 3 | 75075.31 | 2 | 1 | |
| 9999 | 10000 | 792 | 28 | 4 | 130142.79 | 1 | 1 | |

10000 rows × 12 columns

```
In [16]: Newdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   RowNumber             10000 non-null  int64
1   CreditScore           10000 non-null  int64
2   Age                   10000 non-null  int64
3   Tenure                10000 non-null  int64
4   Balance               10000 non-null  float64
5   NumOfProducts         10000 non-null  int64
6   HasCrCard             10000 non-null  int64
7   Geography_France      10000 non-null  uint8
8   Geography_Germany     10000 non-null  uint8
9   Geography_Spain       10000 non-null  uint8
10  Gender_Female         10000 non-null  uint8
11  Gender_Male           10000 non-null  uint8
dtypes: float64(1), int64(6), uint8(5)
memory usage: 595.8 KB
```

```
In [17]: from sklearn.preprocessing import scale
SNewdf=scale(Newdf)
```

```
In [18]: names=Newdf.columns
names
```

```
Out[18]: Index(['RowNumber', 'CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
               'HasCrCard', 'Geography_France', 'Geography_Germany', 'Geography_Spain',
               'Gender_Female', 'Gender_Male'],
              dtype='object')
```

```
In [19]: S_newdf=pd.DataFrame(SNewdf, columns=names)
S_newdf
```

```
Out[19]:
```

| | RowNumber | CreditScore | Age | Tenure | Balance | NumOfProducts | HasCrCard | Geog |
|------|-----------|-------------|-----------|-----------|-----------|---------------|-----------|------|
| 0 | -1.731878 | -0.326221 | 0.293517 | -1.041760 | -1.225848 | -0.911583 | 0.646092 | |
| 1 | -1.731531 | -0.440036 | 0.198164 | -1.387538 | 0.117350 | -0.911583 | -1.547768 | |
| 2 | -1.731185 | -1.536794 | 0.293517 | 1.032908 | 1.333053 | 2.527057 | 0.646092 | |
| 3 | -1.730838 | 0.501521 | 0.007457 | -1.387538 | -1.225848 | 0.807737 | -1.547768 | |
| 4 | -1.730492 | 2.063884 | 0.388871 | -1.041760 | 0.785728 | -0.911583 | 0.646092 | |
| ... | ... | ... | ... | ... | ... | ... | ... | |
| 9995 | 1.730492 | 1.246488 | 0.007457 | -0.004426 | -1.225848 | 0.807737 | 0.646092 | |
| 9996 | 1.730838 | -1.391939 | -0.373958 | 1.724464 | -0.306379 | -0.911583 | 0.646092 | |
| 9997 | 1.731185 | 0.604988 | -0.278604 | 0.687130 | -1.225848 | -0.911583 | -1.547768 | |
| 9998 | 1.731531 | 1.256835 | 0.293517 | -0.695982 | -0.022608 | 0.807737 | 0.646092 | |
| 9999 | 1.731878 | 1.463771 | -1.041433 | -0.350204 | 0.859965 | -0.911583 | 0.646092 | |

10000 rows × 12 columns



```
In [47]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(S_newdf, Odf, test_size=0.2)
```

```
In [48]: # !pip install xgboost
from xgboost import XGBClassifier
xgb=XGBClassifier()
```

```
In [49]: xgb.fit(x_train,y_train)
```

```
[11:18:17] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.5.1/src/learner.cc:1115: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
```

```
Out[49]: XGBClassifier(base_score=0.5, booster='gbtree', colsample_bylevel=1, colsample_bynode=1, colsample_bytree=1, enable_categorical=False, gamma=0, gpu_id=-1, importance_type=None, interaction_constraints='', learning_rate=0.300000012, max_delta_step=0, max_depth=6, min_child_weight=1, missing=nan, monotone_constraints='()', n_estimators=100, n_jobs=8, num_parallel_tree=1, predictor='auto', random_state=0, reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1, tree_method='exact', validate_parameters=1, verbosity=None)
```

```
In [50]: y_pred=xgb.predict(x_test)
```

```
In [51]: y_pred
```

```
Out[51]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [52]: from sklearn.metrics import accuracy_score
```

```
In [54]: Accuracy_Score=accuracy_score(y_test, y_pred)
Accuracy_Score
```

```
Out[54]: 0.85
```

```
In [ ]:
```